

Lecture 19 Intractability I: P, NP, and NP-completeness

L19.1
6.046
05/02/2017

Today

- Problem types: decision, search, optimization
- Polynomial time (P) vs. Non-deterministic polynomial time (NP)
- Reductions
- NP-completeness
- Cook's theorem

Reading for this and the next lecture:
CLRS Chpt 34

Problems: Seemingly related problems can be of vastly different difficulties

Example 1: Shortest vs. Longest Path

Shortest path: Given a weighted graph $G = (V, E, w)$ and vertices $s \neq t$, find the shortest (simple) path from s to t . \rightarrow Polynomial-time algs. exist

Longest path: Given the same input, find the longest (simple) path from s to t .
 \rightarrow No polynomial-time alg. known

Example 2: MST vs. TSP

L19.2

Minimum Spanning Tree (MST): Given a weighted graph $G = (V, E, w)$, find a spanning tree of minimum weight. \rightarrow Polynomial-time algs. exist: $O(|E| \log |V|)$

Traveling Salesman Problem (TSP): Given the same input, find a spanning simple cycle of minimum weight. \rightarrow No polynomial-time alg. known: Current best $O(|V| \cdot 2^{|V|})$ ops

Example 3: Linear vs. Integer Programming

Linear Programming: Given a matrix A and vectors \vec{b} & \vec{c} , find the real vector \vec{x} that maximizes $\vec{c} \cdot \vec{x}$ subject to the constraints $A\vec{x} \leq \vec{b}$.
 \rightarrow Polynomial-time algs. exist.

Integer Programming: Given a matrix A and vectors \vec{b} & \vec{c} , find the integer vector \vec{x} that maximizes $\vec{c} \cdot \vec{x}$ subject to the constraints $A\vec{x} \leq \vec{b}$.
 \rightarrow No polynomial-time alg. known.

Example 4: 2d- vs. 3d-Matching

L19.3

2d-Matching: Given a bipartite graph connecting men to compatible women, find a perfect matching of men to women. → Poly-time algs. exist.

3d-Matching: Given a tri-partite graph connecting men to compatible women, men to compatible restaurants, and women to compatible restaurants, find a collection of man-woman-restaurant triplets so that each man, woman, and restaurant is in a triplet and none are in more than one triplet.

→ No polynomial-time algorithm known.

Conclusions:

(1) Seemingly small changes to problem statements can lead to drastic changes in our algorithmic understanding (and our current ability to solve).

(2) Longest path, TSP, & Integer programming are problems that have been studied intensely by thousands of researchers over a substantial amount of time, yet they have resisted discovery of polynomial-time algorithms.

Questions:

L19.4

- (1) Are there polynomial-time algorithms for these problems waiting to be found?
(→ huge open problem)
- (2) Given a problem, can we tell that it is hard, to avoid "wasting time" trying to find a polynomial time algorithm? (→ yes)

Decision, Search, & Optimization Problems

MST (**optimization** version): Given a weighted graph $G = (V, E, w)$, find a spanning tree of minimum weight. Result is a tree or "graph not connected".

MST (**search** version): Given a weighted graph $G = (V, E, w)$ and a budget K , find a spanning tree whose weight $\leq K$ or report that none exists.
Result is a tree or "insufficient budget" or "graph not connected".

MST (**decision** version): Given a weighted graph $G = (V, E, w)$ and a budget K , decide whether or not there exists a spanning tree with weight $\leq K$.
Result is "yes" or "no".

The existence of a polynomial-time solution to the optimization version of the problem implies a polynomial-time solution to the search version (Why?).

L19.5

Likewise, the existence of a polynomial-time solution to the search version implies one for the decision version.

[Note that sometimes the inference goes in the other direction. For example, a poly-time solution of a search problem can lead to a poly-time solution to the corresponding optimization version by carrying out binary search to find the optimal objective value.]

Conversely, an inability to solve the decision version in poly-time \Rightarrow an inability to solve the search version in poly-time \Rightarrow an inability to solve the optimization version in poly-time. (Why? $P \Rightarrow Q \Leftrightarrow \neg Q \Rightarrow \neg P$)

↳ For showing intractability, we tend to focus on decision version.

Thus, we will be concerned with the decision versions of our problems:

L19.6

Shortest Path: Is there a path of length at most K from s to t?

Longest Path: Is there a simple path of length at least K from s to t?

MST: Is there a spanning tree of weight $\leq K$?

TSP: Is there a spanning simple cycle of weight $\leq K$?

Linear Programming: Is there a real vector \vec{x} s.t. $\vec{c} \cdot \vec{x} \geq k$ and $A \vec{x} \leq \vec{b}$?

Integer Programming: Is there an integer vector \vec{x} s.t. $\vec{c} \cdot \vec{x} \geq k$ and $A \vec{x} \leq \vec{b}$?

2d-matching: Is there a perfect 2d-matching?

3d-matching: Is there a perfect 3d-matching?

Polynomial vs. Non-deterministic Polynomial Time

A decision problem Π is solvable in polynomial time if there exists a polynomial-time algorithm A such that:

$\forall x: (x \text{ is a "yes" input for } \Pi) \Leftrightarrow (A(x) \text{ outputs "yes"})$

We write $\Pi \in P$ problems solvable in polynomial time.
(shortest path, MST, LP) $\in P$

What can we say for longest path, TSP,
integer programming, and 3d-matching?

L19.7

For every "yes" instance there exists a
polynomially short certificate that can be
verified in polynomial time to establish that
the instance is truly a "yes" instance.

3d-matching: If G is a "yes" instance of a
3d-matching problem, then a certificate for it is a
collection of $\frac{|\mathcal{V}|}{3}$ woman-man-restaurant triplets.
And it can be checked in polynomial time whether
or not the collection is truly a perfect 3d-matching.

Longest path: If we have a "yes" instance of a
longest path problem (G, K, s, t) , then a certificate
for it is a path from s to t whose length is at
least K . Again, it can be checked in polynomial
time whether or not the path is valid and its
length is at least K .

Non-Deterministic Polynomial Time (NP)
captures problems with polynomially
short & polynomial-time verifiable
certificates of their "yes" instances.

L19.8

Formally, we say that $\Pi \in NP$

if there exists a polynomial-time verification algorithm V_Π and a constant C such that

$$\forall x : (\Pi(x) = \text{"yes"}) \Leftrightarrow (\exists \text{certificate } y \text{ s.t. } |y| \leq |x|^C \text{ and } V_\Pi(x, y) = \text{"yes"})$$

Note: (i) NP does not stand for "non-polynomial" but "non-deterministic polynomial" (because these problems can be solved by a non-deterministic polynomial algorithm that guesses the solution and then checks it).

(ii) $P \subseteq NP$

Proof: If $\Pi \in P$, \exists poly-time algorithm A s.t.

$\forall x : (\Pi(x) = \text{"yes"}) \Leftrightarrow (A(x) \text{ outputs "yes"})$. Now define $V_\Pi(x, y) = A(x)$. It follows that $\forall x : (\Pi(x) = \text{"yes"}) \Leftrightarrow (V_\Pi(x, y) = \text{"yes"})$.

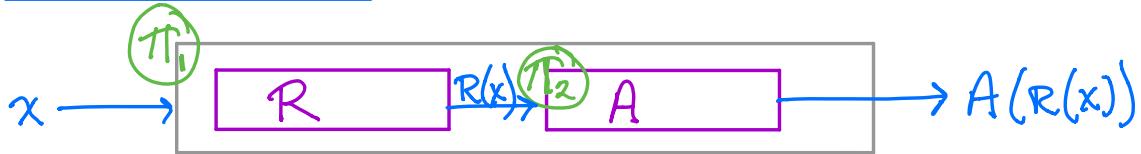
(iii) Believed State of the World



The "P vs NP" problem is considered one of the most important problems in mathematics.

The Clay Institute offers a \$1M prize for its resolution.

Reductions



A polynomial-time reduction from problem Π_1 to problem Π_2 is useful in 2 situations:

- (1) We have a poly-time algorithm for Π_2 and want to obtain a poly-time algorithm for Π_1 .
- (2) We believe Π_1 is a hard problem and want to deduce that Π_2 is also hard.

Formally, a polynomial-time reduction from Π_1 to Π_2 is a polynomial-time algorithm R such that

- if x is an input to Π_1 , then $R(x)$ is an input to Π_2
- $\Pi_1(x) = \text{"yes"} \iff \Pi_2(R(x)) = \text{"yes"}$

And if a poly-time reduction from Π_1 to Π_2 exists, we write $\Pi_1 \leq_p \Pi_2$ (Π_2 is at least as hard as Π_1)

- If $\Pi_1 \leq_p \Pi_2$ and $A(\cdot)$ is a poly-time alg. for Π_2 , then $A(R(\cdot))$ is a poly-time alg. for Π_1 .
- If there is no poly-time alg. for Π_1 and $\Pi_1 \leq_p \Pi_2$, then there is no poly-time algorithm for Π_2 .

Exercise: $\Pi_1 \leq_p \Pi_2$ and $\Pi_2 \leq_p \Pi_3 \Rightarrow \Pi_1 \leq_p \Pi_3$

NP-Completeness

Defn: A problem Π is NP-hard if,
 $\forall \Pi' \in NP, \Pi' \leq_p \Pi$

Defn: A problem Π is NP-complete if,
① $\Pi \in NP$, and
② Π is NP-hard

A problem is NP-hard if it is at least as hard as any problem in NP. It is also NP-complete if it is also in NP (as well as being NP-hard).

We know of numerous NP-complete problems:

- Cook [1971], see below
- Karp [1972] provides a list of 21 NP-complete problems
- Since then the list has grown to contain thousands of NP-complete problems.

It is believed that NP-completeness is one of the biggest exports from Computer Science to the other Sciences and Engineering.

Cook's Theorem

L19.11

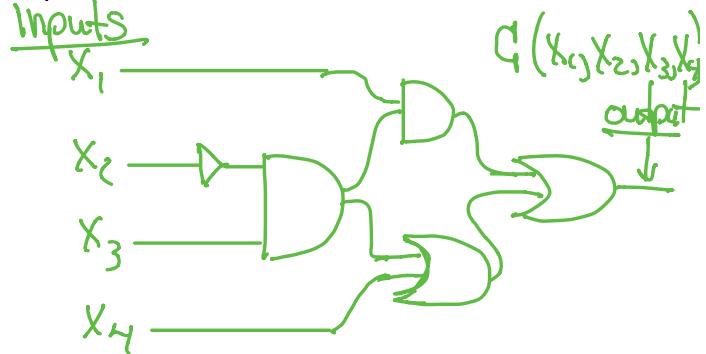
Imagine a circuit made up of 3 types of boolean logic gates: AND, OR, and NOT

fan-in ($= \# \text{ inputs}$): $\underbrace{2}_{\text{unbounded}}$ $\underbrace{1}_{\text{unbounded}}$
fan-out ($= \# \text{ outputs}$): $\text{unbounded} \leftarrow \text{unbounded}$
(Note: CLRS indicates case of unbounded fan-in for AND & OR)
one output with unbounded branching.

Can be represented by graph with nodes labelled

AND, OR, & NOT


Inputs and output are binary variables
 $\{0, 1\}$



Assume no feedback, so graph is DAG, and one output.

Circuit-SAT: Given a circuit $G(x_1, \dots, x_n)$, is there an input for which the output of G is 1?

Theorem [Cook, 1971]: Circuit-SAT is NP-complete.

Proof (some parts sketched):

L19.12

First, we show that Circuit-SAT \in NP. Given a circuit G and a proposed satisfying assignment x_1, \dots, x_n , we can check whether $G(x_1, \dots, x_n) = 1$ in polynomial time.

It remains to show that Circuit-SAT is NP-hard.

- If we pick any problem $\Pi \in$ NP, we want to show that $\Pi \leq_p \text{Circuit-SAT}$.
- Because $\Pi \in$ NP, \exists poly-time V_Π s.t. for every "yes" input x there exists a certificate y s.t. $V_\Pi(x, y) = \text{"yes"}$
- Need to design reduction from Π to Circuit-SAT.
 - (i) Let x be an input to Π .
 - (ii) Reduction builds circuit G_x that is satisfiable iff $\Pi(x) = \text{"yes"}$. $G_x(y)$ is basically an implementation of $V_\Pi(x, y)$.
 - (iii) The technically detailed part of the proof is to show that for any given x , a poly-time alg. $V_\Pi(x, y)$ of y (and of runtime $p(|y|)$) can be implemented by a circuit of size $p_2(|y|)$.

L19.13

Sketch of (cc): On any input y the computation of $V_T(x, y)$ proceeds in $p(|y|)$ steps.

- In every step, the state of the program is a $(|y| + p(|y|) + \log p(|y|))$ -long vector containing input, contents of memory, and program counter.
- The starting state is only the input.
- In going from state t to state $t+1$, the transition is accomplished through a computation that reads parts of the state vector and produces the next state vector. The computation itself can be represented by an AND, OR, NOT circuit.
- Output state: Memory location corresponding to $V_T(x, y)$.
- So the computation of $V_T(x, y)$, for any fixed x and variable y , can be implemented as a layered circuit with $p(|y|)$ layers s.t.
 $C_x(y) = V_T(x, y)$, for all y .

- Thus, C_x is satisfiable iff $\exists y$ s.t. $V_T(x, y) = 1$

