

Admin:

→ Final next Monday ▽

→ Will have practice exams & review sessions. Stay tuned ▽
 ↗ make sure to work through them

Today: Distributed computing

→ Leader election

→ Maximal independent set
 (not "maximum")

So far: Thinking of computation as done on one processor/machine

But: That's not exactly how computation is done today

Computing paradigms:

→ parallel computing (think: multiple processor cores)

(not covered today)

⇒ see 6.816

⇒ how to parallelize a task? (Can it be parallelized?)
 E.g., digging a hole vs. digging a ditch.

→ Can we benefit from having more (identical) workers?

→ distributed computing (think: computer networks/internet)

(today Also: 6.852)

⇒ how to cooperate to solve a joint task?

→ even if some are NOT cooperating

Setup:

- n processors/players
- each has input x_i
- want to compute, for each processor i ,

$$y_i = f_i(x_1, \dots, x_n)$$

E.g. $x_i = \text{numbers}$

$y_i = \text{max of these } n \text{ numbers}$

Crucial: Each f_i might depend on all inputs

⇒ Cooperation essential


(today will ignore situation when some processors are faulty or adversarial)

How to talk?

(today also: G.852)

① Message passing model:


- processors connected in an undirected graph
- in each round: can send/receive messages along edges



→ each proc. know their I/O ports by local name
↑
arbiting

② Shared memory model:

- Processors communicate by reading/writing to shared memory in each round



Think: Msg. board

Important: We assume synchronous model here, i.e., things happen in rounds

Leader election:

"protocol" emphasizes the focus on communication!

(3)

Goal: Run a protocol (algorithm) so as at the end exactly one processor outputs "I am the leader"

Sounds simple but how to do that?

Warning: Impossible if the protocol is deterministic & processors are truly identical

Simplest counterexample:

$n = 2$

Observe:



Round 1: Same starting states
⇒ Same outgoing message
⇒ Same incoming message

Round 2: Same state
⇒ Same outgoing message
⇒ Same incoming message

Round 3: ...

⇒ Either goes on forever or
both processors declare "I am the leader"

⇒ Protocol fails

□

Fundamental problem: Lack of "symmetry breaking" mechanism

Solution I: Make processors non-identical

(4)

\Rightarrow each processor has a unique ID
(email, IP address, MAC, ...)

Simple protocol: (Assume graph is connected)

\rightarrow Each processor i has a local variable \max_i

\rightarrow In each round:

\rightarrow Send \max_i (to all neighbors)

\rightarrow Update $\max_i = \max(\max_i, \{\text{incoming msgs}\})$

\rightarrow After Δ rounds, ($\Delta =$ upper bound on diameter of the graph)
if $\max_i = ID_i$, output
"I am the leader" (otherwise, do nothing)

Solution II: No unique IDs but have randomness

\Rightarrow can use randomness to "manufacture" UIDs!

Protocol:

\rightarrow Choose ID_i at random from the set $\{1, 2, \dots, K\}$ for some $K \geq 1$

\rightarrow Run the protocol for unique ID setting

Clearly: If all ID_i s end up being unique \Rightarrow protocol correct

What is the probability of a collision?

Note: $\forall_{i \neq j} \Pr[ID_i = ID_j] = \frac{1}{K}$

⇒ By union bound over all $\binom{n}{2}$ possible pairs (5)

$$\Pr[\text{collision}] \leq \sum_{i \neq j} \Pr[ID_i = ID_j] = \binom{n}{2} \frac{1}{K} \leq \epsilon$$

if $K \geq \epsilon^{-1} \binom{n}{2}$ for some $\epsilon > 0$

⇒ The protocol succeeds with prob. $\geq 1 - \epsilon$

BUT: Processors do not know if they succeeded!

⇒ Monte Carlo algorithm

Fix: To get a Las Vegas algorithm:

Detect & Repeat (if needed)

⇒ Before "officially" announcing the leader, run a check if there is only one leader
if not, repeat the protocol

⇒ In expectation, need $\leq \frac{1}{(1-\epsilon)}$ repeats

⇒ Expected # of rounds $O\left(\frac{\Delta}{(1-\epsilon)}\right)$

but it is always correct, as needed

checking that can be done deterministically without UIDs, by simple broadcast for Δ rounds

Note: We ignore here time needed for local computations, as we are focused on complexity of reaching a consensus, and this is measured via # of rounds of communication.

Maximal independent set (MIS) problem

⑥

↳ (NOT "maximum")


"maximal" = such that we can't add any new element to it without the need to "swap out" another element

Goal: A protocol such that at the end each processor outputs a Yes/No decision & the yes-decision processors form a maximal independent set


→ independent \Rightarrow no two yes-processors are neighbors

→ maximal \Rightarrow can't add any more yes-processors without violating independence

Note:



maximal & maximum independent set
(size = 3)



maximal but NOT maximum independent set
(size = 2)

Recall: Maximum independent set is NP-hard

But: Maximal independent set (MIS) is in P

(As in distributed computing we ignore local computation time, in principle, we can solve NP-hard problems too! But this would be an "abuse" of the model.)

Simple protocol for MIS:

- Do leader election → add leader to MIS & make its neigh. inactive
- Repeat $\Rightarrow O(n \cdot \Delta)$ rounds

Can we do better?

Luby's (randomized) MIS protocol

Setup: NO UIDs but we have randomness

Protocol:

- All processors are "active" in the beginning
- Protocol proceeds in phases. Each phase = 2 rounds
- In each phase:

Round 1:

- Choose random value $r_i \in \{1, 2, \dots, k\}$ & send to all neighbors
- Receive values from neighbors
- If received values are all $< r_i$, then join the MIS (i.e., output YES)

Round 2:

- If you joined MIS, announce to all neighbors
- If you received such an announcement, decide to NOT join MIS (i.e., output No)
- If you decided YES/No in this phase, become inactive

Observe:

- Final set satisfies independence (since we join MIS only if value of r is uniquely maximal among neighbors; \Rightarrow when this happens, all neighbors output No)

- Final set is maximal (since only way to become inactive is if you or one of your neighbors joins MIS)

Remaining question: How many rounds until done?

(8)

(= all processors inactive)

Lemma: Need (only!) $O(\log n)$ rounds to terminate
(provided $K \gg n^3$, e.g., $K = n^4$ is fine)

In fact: Terminate in $4 \log n$ phases with prob. $\geq 1 - \frac{1}{n}$

Proof: (Will only argue for a line graph)



But: the same ideas + more careful calculations
give the claim for general graphs)

Note: In each phase, if $i \neq j$ then $\Pr[r_i = r_j] = \frac{1}{K} \ll \frac{1}{n^3}$

\Rightarrow union bounding over all pairs & $4 \log n$ first phases

$$\Pr[r_j = r_i, \text{ for some } i \neq j \text{ in phase } l \leq 4 \log n] \leq \frac{4 \log n \cdot \binom{n}{2}}{K} \leq \frac{2n^2 \log n}{K} \ll \frac{1}{n}$$

\Rightarrow wlog can assume all r_i s are always distinct! if $K \geq n^4$

Now: Say an edge (u, v) is active iff both u & v are still active

Key claim: For any edge (u, v) & any phase in which (u, v) starts as active

$$\Pr[(u, v) \text{ becomes inactive}] \geq \frac{1}{2}$$

Proof: Consider cases:

①  (no active edges incident to (u, v))

\Rightarrow either $r_u > r_v$ or vice versa

(with prob. $\geq \frac{1}{2}$)

\Rightarrow either u or v output Yes \Rightarrow (u, v) becomes inactive

② (one active edge incident to (u,v)) ⑨

\Rightarrow with prob $\frac{1}{2}$, $r_u > r_v \Rightarrow u$ outputs Yes

③ (two active edges incident to (u,v)) $\Rightarrow (u,v)$ becomes inactive with prob $\geq \frac{1}{2}$

$\Rightarrow \left. \begin{aligned} \Pr[r_u < r_v \& r_u < r_v] &= \frac{1}{4} \\ \Pr[r_u > r_v \& r_z < r_u] &= \frac{1}{4} \end{aligned} \right\} \text{disjoint events}$

$\Rightarrow \Pr[u \text{ or } v \text{ is a local max}] = \frac{1}{2} \Rightarrow (u,v) \text{ becomes } \underline{\text{inactive}}$
or prob. $\geq \frac{1}{2}$

13

Using the key claim:

$\forall (u,v) \Pr[(u,v) \text{ is still active after } t \text{ phases}] \leq \left(\frac{1}{2}\right)^t$

\Rightarrow union bounding over all $|E|$ edges:

$\Pr[\text{protocol did not terminate after } t \text{ phases}] =$

$= \Pr[\text{Some } (u,v) \text{ still active after } t \text{ phases}] \leq$

$\leq |E| \left(\frac{1}{2}\right)^t \leq \frac{n}{2^t} \leq \frac{1}{n^3} < \frac{1}{n} \text{ if } t \geq 4 \log n$

14